

✓ Optical Character Recognition (OCR) Model: Using TensorFlow and Keras API

Workshop by LU HOU YANG

[TensorFlow Documentation](#)

✓ Install TensorFlow & Configure Environment Locally

Youtube Tutorial & Steps

1. [Youtube](#)
2. [Steps From Youtube Video](#)

-
1. Install [Anaconda](#)
 2. Check [System Requirement](#) for TensorFlow
 3. Check [Software Requirement](#) for TensorFlow
 4. Build Configuration CPU
 - [TF CPU Configuration](#)
 - Create [Conda Environment](#)
 - Install [dependencies](#)

```
# Anything above 2.10 is not supported on the GPU on Windows Native
python -m pip install "tensorflow<2.11"

# Verify the installation:
python -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

5. Build Configuration GPU
 - Check if you GPU supports CUDA [here](#)
 - [TF GPU Configuration](#)
 - [Nvidia CUDA toolkit](#)
 - [cuDNN Archive](#)
 - Create [Conda Environment](#)
 - Install [dependencies](#). Change CUDA toolkit, cuDNN version

```
conda install -c conda-forge cudatoolkit=11.2 cudnn=8.1.0

# Anything above 2.10 is not supported on the GPU on Windows Native
python -m pip install "tensorflow<2.11"

# Verify the installation:
python -c "import tensorflow as tf; print(tf.config.list_physical_devices('GPU'))"
```

✓ 1. Data Collection & Preparation

✓ Data Sources

[Kaggle](#)

[Hugging Face](#)

[NIST](#)

✓ EMNIST Dataset

Extended Modified NIST ([EMNIST](#)) Dataset

Derived from [NIST Special Database 19](#)

- EMNIST ByClass: 814,255 characters. 62 unbalanced classes.
- EMNIST ByMerge: 814,255 characters. 47 unbalanced classes.
- EMNIST Balanced: 131,600 characters. 47 balanced classes.
- EMNIST Letters: 145,600 characters. 26 balanced classes.
- EMNIST Digits: 280,000 characters. 10 balanced classes.
- EMNIST MNIST: 70,000 characters. 10 balanced classes.

✓ Download & Inspect Data

1. Manual Download - [Kaggle, Balanced EMNIST](#)

- [Train CSV](#)
- [Test CSV](#)
- [Mapping](#)

Then, unzip the files in the root of project

2. From Terminal

- Follow steps from here <https://github.com/otenim/Python-EMNIST-Decoder>

3. From IDE - [NIST SD 19, ByMerge](#)

```
import os
import pathlib

DOWNLOAD_DATA_DIR = 'C:\\Users\\User\\Desktop\\Python\\OCR_Workshop\\data'

download_data, by_merge, unzip

download_data_dir = pathlib.Path(DOWNLOAD_DATA_DIR)
if not data_dir.exists():
    tf.keras.utils.get_file(
        'by_merge.zip',
        origin='https://s3.amazonaws.com/nist-srd/SD19/by_merge.zip',
        extract=True,
        cache_dir='.',
        cache_subdir='data'
    )
```

4. Online Google Colab/Jupyter Notebook

Download train/test CSV files, ASCII mapping

```
!gdown 1ruHQJG2pPaeoSpvexdGgKQrwmZPE-Req
!gdown 1bHJFAdbKr_55HSVsjaauZpI9HG009UsZ
!gdown 1hf0mtpXwECXi_IvgwCQngSzmPe-DH1aT
```

Unzip the files

```
!unzip -q /content/emnist-balanced-train.csv.zip
!unzip -q /content/emnist-balanced-test.csv.zip
```

See mapping

```
!cat /content/emnist-balanced-mapping.txt
```

```
# download
```

```
!gdown 1ruHQJG2pPaeoSpvexdGgKQrwmZPE-Req
```

```
!gdown 1bHJFAdbKr_55HSVsjaauZpI9HG009UsZ
```

```
!gdown 1hf0mtpXwECXi_IvgwCQrgSzmPe-DHlaT
```

```

↳ Downloading...
From (original): https://drive.google.com/uc?id=1ruHQJG2pPaeoSpvexdGgKQrwmZPE-Req
From (redirected): https://drive.google.com/uc?id=1ruHQJG2pPaeoSpvexdGgKQrwmZPE-Req&confirm=t&uuid=a33ad1ee-cb33-4c0b-8878-e56f45108864
To: /content/emnist-balanced-train.csv.zip
100% 39.0M/39.0M [00:01<00:00, 19.7MB/s]
Downloading...
From: https://drive.google.com/uc?id=1bHJFAdbKr\_55HSVsjaauZpI9HG009UsZ
To: /content/emnist-balanced-test.csv.zip
100% 6.53M/6.53M [00:00<00:00, 11.6MB/s]
Downloading...
From: https://drive.google.com/uc?id=1hf0mtpXwECXi\_IvgwCQrgSzmPe-DHlaT
To: /content/emnist-balanced-mapping.txt
100% 326/326 [00:00<00:00, 1.31MB/s]
```

```
# unzip
```

```
!unzip -q /content/emnist-balanced-train.csv.zip
```

```
!unzip -q /content/emnist-balanced-test.csv.zip
```

```
# see mapping
```

```
!cat /content/emnist-balanced-mapping.txt
```

```

↳ 0 48
1 49
2 50
3 51
4 52
5 53
6 54
7 55
8 56
9 57
10 65
11 66
12 67
13 68
14 69
15 70
16 71
17 72
18 73
19 74
20 75
21 76
22 77
23 78
24 79
25 80
26 81
27 82
28 83
29 84
30 85
31 86
32 87
33 88
34 89
35 90
36 97
37 98
38 100
39 101
40 102
41 103
42 104
43 110
44 113
45 114
46 116
```

ASCII chart

Decimal - Binary - Octal - Hex – ASCII Conversion Chart

Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII	Decimal	Binary	Octal	Hex	ASCII
0	00000000	000	00	NUL	32	00100000	040	20	SP	64	01000000	100	40	@	96	01100000	140	60	`
1	00000001	001	01	SOH	33	00100001	041	21	!	65	01000001	101	41	A	97	01100001	141	61	a
2	00000010	002	02	STX	34	00100010	042	22	"	66	01000010	102	42	B	98	01100010	142	62	b
3	00000011	003	03	ETX	35	00100011	043	23	#	67	01000011	103	43	C	99	01100011	143	63	c
4	00000100	004	04	EOT	36	00100100	044	24	\$	68	01000100	104	44	D	100	01100100	144	64	d
5	00000101	005	05	ENQ	37	00100101	045	25	%	69	01000101	105	45	E	101	01100101	145	65	e
6	00000110	006	06	ACK	38	00100110	046	26	&	70	01000110	106	46	F	102	01100110	146	66	f
7	00000111	007	07	BEL	39	00100111	047	27	'	71	01000111	107	47	G	103	01100111	147	67	g
8	00001000	010	08	BS	40	00101000	050	28	(72	01001000	110	48	H	104	01101000	150	68	h
9	00001001	011	09	HT	41	00101001	051	29)	73	01001001	111	49	I	105	01101001	151	69	i
10	00001010	012	0A	LF	42	00101010	052	2A	*	74	01001010	112	4A	J	106	01101010	152	6A	j
11	00001011	013	0B	VT	43	00101011	053	2B	+	75	01001011	113	4B	K	107	01101011	153	6B	k
12	00001100	014	0C	FF	44	00101100	054	2C	,	76	01001100	114	4C	L	108	01101100	154	6C	l
13	00001101	015	0D	CR	45	00101101	055	2D	-	77	01001101	115	4D	M	109	01101101	155	6D	m
14	00001110	016	0E	SO	46	00101110	056	2E	.	78	01001110	116	4E	N	110	01101110	156	6E	n
15	00001111	017	0F	SI	47	00101111	057	2F	/	79	01001111	117	4F	O	111	01101111	157	6F	o
16	00010000	020	10	DLE	48	00110000	060	30	0	80	01010000	120	50	P	112	01110000	160	70	p
17	00010001	021	11	DC1	49	00110001	061	31	1	81	01010001	121	51	Q	113	01110001	161	71	q
18	00010010	022	12	DC2	50	00110010	062	32	2	82	01010010	122	52	R	114	01110010	162	72	r
19	00010011	023	13	DC3	51	00110011	063	33	3	83	01010011	123	53	S	115	01110011	163	73	s
20	00010100	024	14	DC4	52	00110100	064	34	4	84	01010100	124	54	T	116	01110100	164	74	t
21	00010101	025	15	NAK	53	00110101	065	35	5	85	01010101	125	55	U	117	01110101	165	75	u
22	00010110	026	16	SYN	54	00110110	066	36	6	86	01010110	126	56	V	118	01110110	166	76	v
23	00010111	027	17	ETB	55	00110111	067	37	7	87	01010111	127	57	W	119	01110111	167	77	w
24	00011000	030	18	CAN	56	00111000	070	38	8	88	01011000	130	58	X	120	01111000	170	78	x
25	00011001	031	19	EM	57	00111001	071	39	9	89	01011001	131	59	Y	121	01111001	171	79	y
26	00011010	032	1A	SUB	58	00111010	072	3A	:	90	01011010	132	5A	Z	122	01111010	172	7A	z
27	00011011	033	1B	ESC	59	00111011	073	3B	;	91	01011011	133	5B	[123	01111011	173	7B	{
28	00011100	034	1C	FS	60	00111100	074	3C	<	92	01011100	134	5C	\	124	01111100	174	7C	
29	00011101	035	1D	GS	61	00111101	075	3D	=	93	01011101	135	5D]	125	01111101	175	7D	}
30	00011110	036	1E	RS	62	00111110	076	3E	>	94	01011110	136	5E	^	126	01111110	176	7E	~
31	00011111	037	1F	US	63	00111111	077	3F	?	95	01011111	137	5F	_	127	01111111	177	7F	DEL

This work is licensed under the Creative Commons Attribution-ShareAlike License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/>

ASCII Conversion Chart.doc Copyright © 2008, 2012 Donald Weiman 22 March 2012

▼ Prepare Data for TensorFlow

Determine Load Function

pd.read_csv

Read a comma-separated values (csv) file into DataFrame. [pd.read_csv](#)

```
pandas.read_csv(
    filepath_or_buffer
)
```

tf.keras.utils.image_dataset_from_directory

Require directory structure for [tf.keras.utils.image_dataset_from_directory](#):

```
main_directory/
...class_a/
.....a_image_1.jpg
.....a_image_2.jpg
...class_b/
.....b_image_1.jpg
.....b_image_2.jpg
```

▼ 2. Visualize Data & Preprocessing

▼ Load Data & Inspect Shape

```
import gc

import tensorflow as tf
import numpy as np
import pandas as pd

import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline

from keras import layers
from keras.models import Sequential
from IPython.display import display, clear_output

gc.collect()
# add any additional imports after this line -----
```

↩ 4

```
# declare global variables
SEED = 42
tf.random.set_seed(SEED)
np.random.seed(SEED)

train_dataframe = pd.read_csv('emnist-balanced-train.csv')
test_dataframe = pd.read_csv('emnist-balanced-test.csv')
mapping = pd.read_csv('emnist-balanced-mapping.txt', sep=' ', header=None)
```

```
display(train_dataframe.head())
train_dataframe.info()
print(f"train_dataframe shape: {train_dataframe.shape}")
```

↩

	45	0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	...	0.524	0.525	0.526	0.527	0.528
0	36	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	43	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	15	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
4	42	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows x 785 columns

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 112799 entries, 0 to 112798
Columns: 785 entries, 45 to 0.533
dtypes: int64(785)
memory usage: 675.6 MB
```

[Kaggle EMNIST](#)

Format


There are six different splits provided in this dataset and each are provided in two formats:

- Binary (see `emnist_source_files.zip`)
- CSV (combined labels and images)
 - Each row is a separate image
 - 785 columns
 - First column = `class_label` (see `mappings.txt` for class label definitions)
 - Each column after represents one pixel value (784 total for a 28 x 28 image)

```
labels = train_dataframe["45"].values
```

```
plt.figure(figsize=(20, 6))
sns.countplot(x=labels)
```

```
del labels
gc.collect()
```



```
def get_mapping(m):
    class_mapping = {}
    ascii_code = m[1].values
    for i, code in enumerate(ascii_code):
        class_mapping[i] = chr(code)

    return class_mapping
```

```
class_mapping = get_mapping(mapping)
print(f"Class Map: {class_mapping}")
```

```
Class Map: {0: '0', 1: '1', 2: '2', 3: '3', 4: '4', 5: '5', 6: '6', 7: '7', 8: '8', 9: '9', 10: 'A', 11: 'B', 12: 'C', 13: 'D', 14: 'E',
```

Preprocessing

Referring to [EMNSIT research paper](#)

- CSV (combined labels and images)
 - Each row is a separate image
 - 785 columns
 - First column = class_label (see mappings.txt for class label definitions)
 - Each column after represents one pixel value (784 total for a 28 x 28 image)

```
# extract label and training data
y_train = np.array(train_dataframe.iloc[:,0].values)
x_train = np.array(train_dataframe.iloc[:,1:].values)

y_test = np.array(test_dataframe.iloc[:,0].values)
x_test = np.array(test_dataframe.iloc[:,1:].values)

print(f"Extracted Labels: {y_train}\n")
print(f"Train data shape: {x_train.shape}")
print(f"Train data shape: {x_test.shape}\n")

del train_dataframe
del test_dataframe
gc.collect()
```

Extracted Labels: [36 43 15 ... 23 31 8]

Train data shape: (112799, 784)
Train data shape: (18799, 784)

9898

```
# normalize
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255


# Reshape the data to have a single color channel (since EMNIST is grayscale)
# and match the input shape expected by the model
x_train = x_train.reshape((-1, 28, 28, 1))
x_test = x_test.reshape((-1, 28, 28, 1))
print(f"Shape before transpose: {x_train.shape}")
print(f"Mean before transpose: {np.mean(x_train[116])}\n")
```

Shape before transpose: (112799, 28, 28, 1)
Mean before transpose: 0.11454582214355469

✓ Show 5 samples from original data

```
fig, axes = plt.subplots(1, 5, figsize=(12, 12))
```

```
for i, ax in enumerate(axes.flatten()):
    ax.imshow(x_train[i])
    ax.set_title(y_train[i])
```



```
# transpose to reverse effect from converting png to csv (2D -> 1D)
def transpose_data(x):
    return np.transpose(np.squeeze(x))[..., np.newaxis]
```

```
x_train = np.array([transpose_data(x) for x in x_train])
x_test = np.array([transpose_data(x) for x in x_test])
print(f"Shape after transpose: {x_train.shape}")
print(f"Mean after transpose: {np.mean(x_train[116])}\n")
```

Shape after transpose: (112799, 28, 28, 1)
Mean after transpose: 0.11454582214355469

```
# get a sample image for future comparison
sample_image = x_train[116]
sample_label = y_train[116]
print(f"Sample image shape: {sample_image.shape}")
print(f"Sample label: {class_mapping.get(sample_label)}")
```

```
↗ Sample image shape: (28, 28, 1)
Sample label: F
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))
test_dataset = tf.data.Dataset.from_tensor_slices((x_test, y_test))
print(f"train_dataset shape: {train_dataset.element_spec}")
```

```
↗ train_dataset shape: (TensorSpec(shape=(28, 28, 1), dtype=tf.float32, name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))
```

```
del x_train
del x_test
del y_train
del y_test
gc.collect()
```


```
↗ 25271
```

```
train_dataset = train_dataset.cache().shuffle(10000).batch(16).prefetch(tf.data.AUTOTUNE)
test_dataset = test_dataset.cache().batch(16)
```

✓ Show 5 samples from preprocessed data

```
for img_ds, labels_ds in train_dataset.take(1):
    fig, axes = plt.subplots(1, 5, figsize=(12, 12))

    for i, ax in enumerate(axes.flatten()):
        ax.imshow(img_ds[i])
        ax.set_title(class_mapping.get(labels_ds[i].numpy()))
```



✓ 3. Build Model

Neural Network (NN)

Convolutional Neural Network (CNN)

✓ Build Sequential Model


```

input_shape = img_ds.shape[1:]
num_classes = len(class_mapping)

model = Sequential([
    layers.Input(shape=input_shape),
    layers.Conv2D(filters=32, kernel_size=(6, 6), activation='relu', padding='same'),
    layers.MaxPooling2D((3, 3)),
    layers.Conv2D(filters=64, kernel_size=(4, 4), activation='relu', padding='same'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(filters=128, kernel_size=(3, 3), activation='relu'),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_classes)
])

```

✓ 4. Compile, Fit & Save Model

Compile: specifying a loss, metrics, and an optimizer

✓ Optimizers

Adam

Adaptive Moment Estimation is an algorithm for optimization technique for gradient descent. [Read more here](#)

RMSprop

RMSprop (Root Mean Squared Propagation) is an adaptive learning rate optimization algorithm. [Read more here](#)

✓ Loss Functions

Categorical Cross-Entropy

Computes the crossentropy loss between the labels (**one-hot encoded vectors**) and predictions. [Read more here](#)

Sparse Categorical Cross-Entropy

Computes the crossentropy loss between the labels (**integer values**) and predictions. [Read more here](#)

✓ Metrics

Accuracy

Loss

Val Accuracy

Val Loss

✓ Compile

```

model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy']
)

```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	1184
max_pooling2d (MaxPooling2D)	(None, 9, 9, 32)	0
conv2d_1 (Conv2D)	(None, 9, 9, 64)	32832
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0
conv2d_2 (Conv2D)	(None, 2, 2, 128)	73856
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 64)	32832
dense_1 (Dense)	(None, 64)	4160
dropout (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 47)	3055

=====
Total params: 147919 (577.81 KB)
Trainable params: 147919 (577.81 KB)
Non-trainable params: 0 (0.00 Byte)
=====

Fit Model

EPOCHS = 5

```
early_stopping = tf.keras.callbacks.EarlyStopping(  
    monitor='val_accuracy',  
    patience=3,  
    restore_best_weights=True,  
    min_delta=0,  
    mode='max',  
    verbose=0  
)
```

```
history = model.fit(  
    train_dataset,  
    validation_data=test_dataset,  
    epochs=EPOCHS,  
    callbacks=[early_stopping]  
)
```

Epoch 1/5
7050/7050 [=====] - 42s 5ms/step - loss: 0.9275 - accuracy: 0.7156 - val_loss: 0.4725 - val_accuracy: 0.8402
Epoch 2/5
7050/7050 [=====] - 34s 5ms/step - loss: 0.5264 - accuracy: 0.8308 - val_loss: 0.4435 - val_accuracy: 0.8516
Epoch 3/5
7050/7050 [=====] - 36s 5ms/step - loss: 0.4604 - accuracy: 0.8495 - val_loss: 0.4475 - val_accuracy: 0.8569
Epoch 4/5
7050/7050 [=====] - 33s 5ms/step - loss: 0.4265 - accuracy: 0.8593 - val_loss: 0.4141 - val_accuracy: 0.8632
Epoch 5/5
7050/7050 [=====] - 34s 5ms/step - loss: 0.4003 - accuracy: 0.8667 - val_loss: 0.4314 - val_accuracy: 0.8661

Save Model

```
model.save('ocr_model')
```

5. Test Accuracy of Model

Evaluate

Evaluate Function

```
model.evaluate(test_dataset)
```

```
1175/1175 [=====] - 3s 2ms/step - loss: 0.4314 - accuracy: 0.8661
[0.43141597509384155, 0.8661099076271057]
```

6. Export Model

Python Module

```
class OCDModel(tf.Module):
    def __init__(self, model):
        self.model = model
        self.class_mapping = {0: '0', 1: '1', 2: '2', 3: '3', 4: '4', 5: '5', 6: '6', 7: '7', 8: '8', 9: '9',
                               10: 'A', 11: 'B', 12: 'C', 13: 'D', 14: 'E', 15: 'F', 16: 'G', 17: 'H', 18: 'I',
                               19: 'J', 20: 'K', 21: 'L', 22: 'M', 23: 'N', 24: 'O', 25: 'P', 26: 'Q', 27: 'R',
                               28: 'S', 29: 'T', 30: 'U', 31: 'V', 32: 'W', 33: 'X', 34: 'Y', 35: 'Z', 36: 'a',
                               37: 'b', 38: 'd', 39: 'e', 40: 'f', 41: 'g', 42: 'h', 43: 'n', 44: 'q', 45: 'r',
                               46: 't'}

    @tf.function
    def predict(self, data):
        if isinstance(data, tf.float32):
            result = self.model(data, training=False)
        elif isinstance(data, str):
            img = preprocess_image(data, 'preprocessed_input_img.png')
            result = self.model(img, training=False)
        else:
            raise ValueError("Unsurported data type.\nPlease pass preprocessed image using preprocess_image function.\nOr pass path to image file")

        return result

    def __call__(self, data):
        pred = self.predict(data)
        return self.class_mapping.get(np.argmax(pred.numpy())[0])
```

TensorFlow Lite for Flutter

TensorFlow Lite

- TensorFlow Lite is a way to run TensorFlow models on devices locally, supporting mobile, embedded, web, and edge devices. [Read more here](#)
- Convert model to TFLite [Documentation](#)

Start coding or [generate](#) with AI.

7. UI, Input & Preprocessing Pipeline

Preprocessing pipeline

```
!gdown 1-cTOocQVoNPA14u5oxu6eBgIUG5jg4y3
!gdown 12NGQvSBp0hBIbFidPYAjjKhX17ooJGn
```

```
Downloading...
From: https://drive.google.com/uc?id=1-cTOocQVoNPA14u5oxu6eBgIUG5jg4y3
To: /content/random_scale_img_f.png
100% 14.1k/14.1k [00:00<00:00, 36.0MB/s]
Downloading...
From: https://drive.google.com/uc?id=12NGQvSBp0hBIbFidPYAjjKhX17ooJGn
```

To: /content/sample_input.png
 100% 793/793 [00:00<00:00, 4.22MB/s]

```
import cv2
from scipy.ndimage import gaussian_filter

def preprocess_image(input_image_path, output_image_path):
    # read the original image
    image = cv2.imread(input_image_path, cv2.IMREAD_GRAYSCALE)
    if image is None:
        raise ValueError("Image not found or the path is incorrect")

    # invert the colors of the image
    image = 255 - image

    # apply Gaussian filter with  $\sigma = 1$ 
    image = gaussian_filter(image, sigma=1)

    # extract the region around the character
    # find non-zero pixels (characters)
    coords = cv2.findNonZero(image)
    x, y, w, h = cv2.boundingRect(coords)

    # crop the image to the bounding box
    cropped_image = image[y:y+h, x:x+w]

    # center the character in a square image
    # calculate the size of the new image (keeping the aspect ratio)
    max_side = max(w, h)
    square_image = np.zeros((max_side, max_side), dtype=np.uint8)

    # compute the offset to center the character
    x_offset = (max_side - w) // 2
    y_offset = (max_side - h) // 2

    # place the cropped image in the center of the square image
    square_image[y_offset:y_offset+h, x_offset:x_offset+w] = cropped_image

    # add a 2-pixel border
    padded_image = cv2.copyMakeBorder(square_image, 2, 2, 2, 2, cv2.BORDER_CONSTANT, value=0)

    # down-sample to 28x28 using bi-cubic interpolation
    downsampled_image = cv2.resize(padded_image, (28, 28), interpolation=cv2.INTER_CUBIC)

    # scale intensity values to [0, 255]
    # convert image to have values in range [0, 255]
    final_image = cv2.normalize(downsampled_image, None, 0, 255, cv2.NORM_MINMAX)

    # normalize values between [0, 1]
    final_image = final_image/255.0

    # verify that preprocessing is consistent with data
    print(f"Mean input image: {np.mean(final_image)}\n")
    plt.imshow(final_image[...], np.newaxis))

    # save the final processed image
    cv2.imwrite(output_image_path, final_image[...], np.newaxis))

    # add batch shape, and channel
    return final_image[np.newaxis, ..., np.newaxis]


# compare sample from dataset & preprocessed input image
fig, axes = plt.subplots(1, 3, figsize=(12, 8))
axes[0].imshow(sample_image)
axes[0].set_title(np.mean(sample_image));

preprocessed_img_nist = preprocess_image('sample_input.png', 'preprocessed_sample_input.png')
img1 = np.squeeze(preprocessed_img_nist, axis=0)
axes[1].imshow(img1)
axes[1].set_title(np.mean(img1));

preprocessed_img_f = preprocess_image('random_scale_img_f.png', 'preprocessed_random_scale_img_f.png')
img2 = np.squeeze(preprocessed_img_f, axis=0)
axes[2].imshow(img2)
axes[2].set_title(np.mean(img2));
```

Mean input image: 0.10560724289715887

Mean input image: 0.09445278111244497



8. Load & Use Pretrained Model

Run prediction on Google Colab

```
model = tf.keras.models.load_model('ocr_model')
pred = model.predict(preprocessed_img_f)
print(f"Prediction: {class_mapping.get(np.argmax(pred[0]))}")
```

```
plt.figure()
plt.imshow(np.squeeze(preprocessed_img_f, axis=0))
plt.title(class_mapping.get(np.argmax(pred[0])));
```

WARNING:tensorflow:5 out of the last 5 calls to <function Model.make_predict_function.<1
1/1 [=====] - 0s 118ms/step
Prediction: F

